

Guided Walkthrough for rational agent framework

August 10, 2023

This document walks through how to apply the rational agent framework to your visualization experiment. We provide general guidance and walk through the forecast visualization example in the paper in more detail.

1 General guide for applying the rational agent framework

1.1 Identifying the decision problem

Before we get started to applying the framework to a visualization experiment, we first need to identify the decision-making problem that forms the task in the experiment. Identifying the decision problem involves identifying:

- **states:** the description of reality we want help users understand with visualization, e.g. whether the temperature is above or below freezing in our weather forecast example.
- **data generating model:** the distribution of the distributions that states are generated from. In the weather forecast example, the data generating model is a uniform distribution on $\{Normal(5, 2), Normal(5, 3), Normal(5, 4), Normal(5, 5)\}$.
- an **action space:** the report space that the experiment requires users to use, e.g. whether or not to salt the parking lot.
- **signals:** the visualizations you show to users to aid their decisions.
- a **payoff** or **scoring rule:** the payoff function used in experiment to incentivize/score users' behaviors. In general, you can think of the scoring rule in the following two forms.
 - payoff: $A \times \Theta \rightarrow \mathbb{R}$, a function mapping an action and the state to a real-valued payoff. e.g. how many points you lose if you salt when it doesn't freeze or don't salt but it freezes.

- proper score: $\Delta(\Theta) \times \Theta \rightarrow \mathbb{R}$ a function mapping the probability distribution over state and the state to a real-valued payoff. The proper score applies the optimal action given the belief (a probability distribution), and is a proper measure of the quality of a decision. See Gneiting and Raftery, "Strictly Proper Scoring Rules, Prediction, and Estimation" (2007) for more information.

There is a one-to-one mapping between a payoff and a proper score. We can always construct a proper score from a payoff.

Note that the use of scoring rules may have two purposes:

- to incentivize behavioral agents in an experiment, and
- to evaluate the effect of seeing a visualization in an experiment.

An experimenter may not use the same scoring rule for incentives and for evaluation. For example, the experimenter may use a flat payment to reward behavioral agents. Hence, the pre-experimental analysis has two purposes with different use of scoring rules:

- if it is believed that behavioral agents are incentivized by monetary rewards, the experimenter should pre-experimentally calculate the value of information in real money, and evaluate the incentive.
- The experimenter can also use a proper scoring rule solely for the purpose of evaluation. The rational baseline, benchmark and value of information can be calculated pre-experimentally to quantify the “room for improvement” in the experiment, thus to quantify information asymmetry.

In either case, the scoring rule is taken as input to our framework. The experimenter can select any payoff function, and construct a proper score, depending on the assumptions they make about the participants.

1.2 Pre-experiment analysis

In this section, we will show how to calculate the three pre-experiment quantities (rational baseline, rational benchmark, and value of information) that the rational agent framework offers. The rational agent’s prior belief is known before seeing the visualization, and represents the unconditional probability that each state is drawn according to the data generating model. Given a visualization signal, the rational agent will update their beliefs about the state based on the joint distribution between visualizations and states. We provide the pseudocode in Alg. 1 and Alg. 2 to calculate the pre-experiment quantities.

For the rational baseline, we assume that the rational agent has no access to the signal, i.e. responds based on only the prior belief. Therefore, when simulating this agent on real experiment data (Alg. 1), the agent won’t care about the visualization or any provided information but always take the same action based on the prior belief.

Algorithm 1: Rational baseline

Input: the experimental data D with each row representing one experimental trial, prior beliefs about the state $p(d)$, the action space A , and the scoring rule S

Output: the rational baseline R_\emptyset

```
belief  $\leftarrow E_{d \sim p(d)}(d)$ ;  
/* the prior belief about states */  
action  $\leftarrow \arg \max_{a \sim A} E_{\theta \sim \text{belief}}(S(a, \theta))$ ;  
/* the action made on prior belief */  
payoff  $\leftarrow 0$ ;  
for row  $\in D$  do  
    | dist  $\leftarrow$  the actual distribution that states are drawn used in row;  
    | payoff  $\leftarrow$  payoff +  $E_{\theta \sim \text{dist}}(S(\text{action}, \theta))$ ;  
end  
 $R_\emptyset \leftarrow \text{payoff} / \# \text{ of row in } D$ ;  
/* rational baseline is the rational agent's expected payoff  
   in the experiment under only their prior belief */
```

For the rational benchmark, we assume the rational agent has access to the visualization (signal). After seeing the visualization (signals), the rational agent will update their belief about the state based on the joint distribution between visualizations and states. For example, the distribution the visualization signals v on the distribution of states d is $p(v|d)$ and the prior belief about the states, i.e. the data generating model, is $p(d)$. Then the posterior belief about the state could be $p(d|v) = p(v|d) \cdot p(d)/p(v)$. See the detail algorithm in Alg. 2.

You can calculate the value of information in the rational agent framework as the difference between the rational benchmark and baseline, i.e. the maximum improvement that can be made when agent has the prior knowledge and follows Bayes rule to update their belief.

$$\Delta = R_\emptyset - R_V$$

1.3 Post-experiment analysis

The rational agent framework also offers two post-experiment measures (behavioral score and calibrated behavioral score) to quantify behavioral performance. The behavioral score B is the payoff that the behavioral agents (i.e. real users) get or are expected to get in the experiment, which is always below the rational benchmark R_V and can be either above or below the rational baseline R_\emptyset . Importantly, if the behavioral score is below the rational baseline, then from the scores alone we cannot reject the hypothesis that the behavioral agent got no useful information from the visualization, because even with no information, the rational agent performs better. Therefore, we propose the calibrated behavioral score, where you can use the rational agent to calibrate the behavioral

Algorithm 2: Rational benchmark

Input: the experimental data D with each row representing one experimental trial, prior beliefs about the state $p(d)$, the action space A , and the scoring rule S

Output: the rational benchmark R_V

```
payoff  $\leftarrow$  0;
for row  $\in D$  do
     $v \leftarrow$  visualization used in row;
     $p(d|v) = p(v|d) \cdot p(d)/p(v)$ ;
    /* the posterior belief about the distribution of states
       */
     $belief \leftarrow E_{d \sim p(d|v)}(d)$ ;
    /* the posterior belief about states */
     $action \leftarrow \arg \max_{a \in A} E_{\theta \sim belief}(S(a, \theta))$ ;
    /* the action made on posterior belief */
     $dist \leftarrow$  the actual distribution of states used in row;
     $payoff \leftarrow payoff + E_{\theta \sim dist}(S(action, \theta))$ ;
    /* the expected payoff with posterior information */
end
 $R_V \leftarrow payoff / \#$  of row in  $D$ ;
```

score to fall within the interval between the rational baseline and the rational benchmark, to understand how much information the behavioral agents didn't get.

You can follow Alg. 3 to estimate the payoff of behavioral agents' actions.

Algorithm 3: Behavioral score

Input: the experimental data D with each row representing one experimental trial, and the scoring rule S

Output: the behavioral scores *behavioral*

```
payoff  $\leftarrow$  0;
for row  $\in D$  do
     $dist \leftarrow$  the actual distribution of states used in row;
     $action \leftarrow$  action made by the behavioral agent in row;
     $payoff \leftarrow payoff + E_{\theta \sim dist}(S(action, \theta))$ ;
    /* the expected score of behavioral action */
end
 $behavioral \leftarrow payoff / \#$  of row in  $D$ ;
```

For the calibrated behavioral score, you simulate the rational agent as in calculating the rational baseline and the rational benchmark, but the rational agent is not using the visualizations as signals but rather using the actions made by behavioral agents. This means that when seeing a behavioral agents'

action a , the rational agent will update the belief about states based on the joint distribution between behavioral actions and states, i.e. $p(d|a) = p(a|d) \cdot p(d)/p(a)$ (Alg. 4).

Algorithm 4: Calibrated behavioral score

Input: the experimental data D with each row representing one experimental trial, prior beliefs about state $p(d)$, the action space A , and the scoring rule S

Output: the calibrated behavioral score $benchmark$

```

payoff  $\leftarrow$  0;
for row  $\in D$  do
    signal  $\leftarrow$  action made by the behavioral agent in row
     $p(d|signal) = p(signal|d) \cdot p(d)/p(signal)$ ;
    /* the posterior belief of the distribution of states in
       calibration */
    belief  $\leftarrow E_{d \sim p(d|signal)}(d)$ ;
    /* the belief of the states */
    action  $\leftarrow \arg \max_{a \in A} E_{\theta \sim belief}(S(a, \theta))$ ;
    /* the action made on posterior belief */
    dist  $\leftarrow$  the actual distribution of states used in row;
    payoff  $\leftarrow$  payoff +  $E_{\theta \sim dist}(S(action, \theta))$ ;
    /* the expected payoff with posterior information */
end
benchmark  $\leftarrow$  payoff / # of row in  $D$ ;
```

2 Walkthrough example: weather forecast visualization

2.1 Setting up for R code

```

> library(dplyr)
> library(tidyr)
> library(modelr)
> knitr::opts_chunk$set(echo = TRUE)
```

2.2 Identifying the decision problem

We will assume an experiment that is a slight variation on the weather forecast example in the paper, where we are asking users to report two actions:

- binary decision: salt the parking lot / not salt
- belief about the possibility of freezing.

The decision-making problem of salting / not salting can be formalized as:

- states: $\theta \in \Theta = \{0 = \text{not freezing}, 1 = \text{freezing}\}$.
- Assume daily low temperature $t \sim N(\mu, \sigma^2)$; $\Pr[\theta = 1] = \Pr[t \leq 0]$; $\mu = 5$ fixed and σ uniformly from $\{2, 3, 4, 5\}$. The information provided by different visualization strategies:
 - mean: the mean display is not more informative than prior, just knowing $\mu = 5$.
 - CI, gradient, HOPs: seeing the visualization is the same as seeing the distribution of daily low temperature, or equivalently knowing the standard deviation σ of temperature.
- action space: $a \in A = \{0 = \text{not salt}, 1 = \text{salt}\}$
- signals: $v \in V$ showing daily low temperature t .
- scoring rule:
 - payoff: $S_{decision}(a, \theta)$ for reporting the decision of salting, see equation 2 in the paper.
 - proper score: $S_{belief}(b, \theta)$ for reporting the belief of freezing, see eq. 1 below.

$$S_{belief}(b, \theta) = S_{decision}(\arg \max_{a \in A} S_{decision}(a, b), \theta) \quad (1)$$

The following R code defines the scoring rule $S_{decision}$ and S_{belief} . $S_{decision}$ scores the binary action that agents take, where we follow the payoff function in equation 2 in the paper.

```
> belief_space = seq(0, 1, 0.02)
> decision_space = c(0, 1)
> decision_scoring_rule = function(decision, state) {
+   # See equation 3 in the paper
+   -100 * (1 - decision) * state - 10 * decision * (1 - state)
+ }
```

We construct a proper score S_{belief} scoring the belief. S_{belief} is constructed from the payoff $S_{decision}$ by applying the optimal binary decision to the belief, i.e. assuming the agent makes the decision rationally on their belief.

```
> belief_scoring_rule = function(belief, state) {
+   # apply the optimal decision based on the belief
+   decision = decision_space[apply(matrix(unlist(lapply(decision_space, function(d) {
+     decision_scoring_rule(d, belief)
+   })), ncol = length(belief)), 2, which.max)]
+   decision_scoring_rule(decision, state)
+ }
```

2.3 For the purposes of this walkthrough only: Simulate behavioral data using quantal response equilibrium

For the purposes of our walkthrough of the rational agent approach, we generate mock experimental data for the weather forecast example described above using quantal response equilibrium¹ (a type of behavioral model that assumes the probability of any particular action being chosen is positively related to the payoff from that action, i.e. very bad actions are unlikely), where we keep the distribution of sigma, i.e. the data generating model, as uniform distribution.

```
> p_d = c(0.25, 0.25, 0.25, 0.25)
> # p(d) the prior belief about states
>
> sigma_choices = c(2, 3, 4, 5)
> mu = 5
> data = data.frame(
+   mu = mu,
+   sigma = sigma_choices,
+   vis = c("mean", "mean+interval", "gradient", "HOPs")
+ ) %>%
+   data_grid(mu, sigma, vis) %>%
+   rowwise() %>%
+   mutate(
+     freezing_prob = pnorm(0, mu, sigma),
+     behavioral_decision = list(sample(decision_space, 100,
+       replace = TRUE,
+       # quantal response equilibrium
+       prob = exp(0.2 * decision_scoring_rule(decision_space, freezing_prob)))),
+     behavioral_belief = list(sample(belief_space, 100,
+       replace = TRUE,
+       # quantal response equilibrium
+       prob = exp(belief_scoring_rule(belief_space, freezing_prob))))
+   ) %>%
+   unnest(cols = c(behavioral_decision, behavioral_belief))
> head(data)
```

A tibble: 6 × 6

	mu	sigma	vis	freezing_prob	behavioral_decision	behavioral_belief
	<dbl>	<dbl>	<chr>	<dbl>	<dbl>	<dbl>
1	5	2	gradient	0.00621	0	0.38
2	5	2	gradient	0.00621	0	0.08
3	5	2	gradient	0.00621	0	0.24
4	5	2	gradient	0.00621	1	0.32
5	5	2	gradient	0.00621	0	0.48
6	5	2	gradient	0.00621	0	0.36

¹<https://www.sciencedirect.com/science/article/pii/S0899825685710238>

2.4 Pre-experiment analysis

For the purpose of walking through the steps entailed in applying the framework, we first show how to calculate the pre-experiment quantities for the weather forecast example by explicit calculations and then walk through it with full code that assumes as behavioral data is the simulated data set generated above.

2.4.1 Explicit Calculations for Weather Forecasting Experiment

Prior We have the prior probability of freezing

$$\begin{aligned} p &= \sum_{\sigma} \Pr[\theta = 1, \sigma] = 0.00155 + 0.01195 + 0.0264 + 0.0397 \\ &= 0.0796. \end{aligned}$$

Posterior The posterior probabilities are $\Pr[\theta = 1|\sigma] = 0.62\%, 4.78\%, 10.56\%, 15.87\%$, relatively for $\sigma = 2, 3, 4, 5$. Here we take $\sigma = 2$ as an example:

$$\Pr[\theta = 1|\sigma = 2] = \Pr_{x \sim N(5, 2^2)}[x \leq 0] = 0.0062.$$

The rational agent framework gives the following quantities:

rational baseline: $R_{\emptyset} = -7.96$.

The prior $p = 0.08$ is optimized at no-salt and gives an expected payoff of -7.96 . The calculation is as follows:

$$\begin{aligned} R_{\emptyset} &= \Pr[\theta = 0] \cdot S(a = 0, \theta = 0) + \Pr[\theta = 1] \cdot S(a = 0, \theta = 1) \\ &= (1 - 0.0796) \times 0 + 0.0796 \times (-100) = -7.96 \end{aligned}$$

visualization optimal: $R_V^{\text{CI}} = R_V^{\text{gradient}} = R_V^{\text{HOPs}} = -5.69$; $R_V^{\text{mean}} = -7.96$.

In CI, gradient, and HOPs, each signal arises with probability $1/4$ and the average of the optimal actions under the induced posteriors gives $R_V = -5.69$. For the visualization of the mean, the rational agent has only the prior information and obtains $R_V^{\text{mean}} = R_{\emptyset} = -7.59$.

The calculation of R_V for CI, gradient, and HOPs is the following:

$$\begin{aligned} R_V^{\text{CI}} &= \sum_{\sigma, \theta} \Pr[\sigma, \theta] \cdot S(\mathbf{1}_{\sigma \in \{4, 5\}}, \theta) \\ &= 0.24845 \times 0 + 0.00155 \times (-100) + 0.23805 \times 0 \\ &\quad + 0.01195 \times (-100) + 0.2236 \times (-10) + 0.0264 \times 0 \\ &\quad + 0.2103 \times (-10) + 0.0397 \times 0 \\ &= -5.69 \end{aligned}$$

rational benchmark: $R_I = \max_{\text{vis}} R_V^{\text{vis}} = -5.69$, the best achievable across visualizations.

value of information: $\Delta = R_I - R_{\emptyset} = 2.27$.

2.4.2 Rational baseline

To calculate the rational baseline, we first calculate the prior belief of rational agent and their associated action choice given these beliefs. Then we calculate what score the rational agent would get under the scoring rule if they complete the experiment with only prior knowledge.

```
> prior_belief = Reduce("+",
+   sapply(sigma_choices, function(m) {pnorm(0, mu, m)})) / length(sigma_choices)
> decision_payoffs = lapply(decision_space, function(a) {
+   decision_scoring_rule(a, prior_belief)
+ })
> prior_reporting_decision = decision_space[[which.max(decision_payoffs)]]
> prior_reporting_decision

[1] 0

> belief_payoffs = lapply(belief_space, function(b) {
+   belief_scoring_rule(b, prior_belief)
+ })
> prior_reporting_belief = belief_space[[which.max(belief_payoffs)]]
> prior_reporting_belief

[1] 0

> # simulate the action of rational agent with only prior knowledge on all experiments
> for (i in 1:nrow(data)) {
+   data[i, "prior_decision_payoff"] =
+     decision_scoring_rule(prior_reporting_decision, data[i, "freezing_prob"])
+   data[i, "prior_belief_payoff"] =
+     belief_scoring_rule(prior_reporting_belief, data[i, "freezing_prob"])
+ }
> prior_payoff = data %>%
+   dplyr::group_by(vis) %>%
+   summarise(prior_decision_payoff = mean(prior_decision_payoff),
+             prior_belief_payoff = mean(prior_belief_payoff))
> prior_payoff

# A tibble: 4 × 3
  vis          prior_decision_payoff prior_belief_payoff
<chr>                <dbl>                <dbl>
1 HOPs                 -7.96                 -7.96
2 gradient             -7.96                 -7.96
3 mean                 -7.96                 -7.96
4 mean+interval        -7.96                 -7.96
```

The result says the optimal expected payoff of the rational agent with only prior knowledge is -7.96 for both reporting decision and belief, and the optimal decision and belief are both 0, meaning that the rational agent with only prior belief would always report 0 and the rational baseline is -7.96 .

2.4.3 Rational benchmark

In the weather forecast example, the visualizations provide perfect information about the variance of data except "mean" visualization, so the updated belief on the distribution of states will converge to a situation that only one distribution of state has possibility with almost 1 and others are all almost 0. In this case, we assume that the rational agent would exactly know which distribution of states is using after viewing one visualization. Formally, assuming that v provides perfect information about the distribution of states d_1 , then

$$\begin{aligned}p(v|d_1) &\approx 1 \\p(v|d_2) &\approx 0 \\&\dots \\p(v|d_n) &\approx 0\end{aligned}$$

so $p(d_1|v) = p(v|d) \cdot p(d)/p(v) \approx 1$ when the visualization is fixed.

```
> for (i in 1:nrow(data)) {
+   # rational agent can infer the exact distribution of states based on signal
+   if (data[[i, "vis"]] == "mean") {
+     posterior_belief = prior_belief
+   } else {
+     posterior_belief = data[[i, "freezing_prob"]]
+   }
+   posterior_reporting_decision = decision_space[[
+     which.max(lapply(decision_space, function(a) {
+       decision_scoring_rule(a, posterior_belief)
+     }))]
+   ]
+   posterior_reporting_belief = belief_space[[
+     which.max(lapply(belief_space, function(b) {
+       belief_scoring_rule(b, posterior_belief)
+     }))]
+   ]
+   data[[i, "posterior_decision_payoff"]] =
+     decision_scoring_rule(posterior_reporting_decision, data[[i, "freezing_prob"]])
+   data[[i, "posterior_belief_payoff"]] =
+     belief_scoring_rule(posterior_reporting_belief, data[[i, "freezing_prob"]])
+ }
> posterior_payoff = data %>%
+   dplyr::group_by(vis) %>%
+   summarise(posterior_decision_payoff = mean(posterior_decision_payoff),
+             posterior_belief_payoff = mean(posterior_belief_payoff))
> posterior_payoff
```

```
# A tibble: 4 × 3
  vis                posterior_decision_payoff posterior_belief_payoff
<chr>                <dbl>                <dbl>
1 HOPs                -5.69                -5.69
2 gradient            -5.69                -5.69
3 mean                -7.96                -7.96
4 mean+interval       -5.69                -5.69
```

When seeing "mean" visualization, the rational agent didn't get the information about sigma, so they can't update the belief, and the rational benchmark for both reporting action and belief turns out to be -7.96 , which is the same as rational baseline. For the other visualizations, the rational benchmark for both reporting action and belief is -5.69 .

2.4.4 Value of information

In the example of weather forecast, the value of information is `posterior_payoff - prior_payoff = 2.27` expect 0 for "mean" visualization.

2.5 Post-experiment analysis

In this section, we will calculate the score of behavioral agents and the score of the calibrated behavioral agents.

2.5.1 Behavioral score

We use the scoring rule on behavioral agents' actions.

```
> for (i in 1:nrow(data)) {
+   data[i, "behavioral_decision_payoff"] =
+     decision_scoring_rule(data[i, "behavioral_decision"], data[i, "freezing_prob"])
+   data[i, "behavioral_belief_payoff"] =
+     belief_scoring_rule(data[i, "behavioral_belief"], data[i, "freezing_prob"])
+ }
> behavioral_payoff = data %>%
+   dplyr::group_by(vis) %>%
+   summarise(behavioral_decision_payoff = mean(behavioral_decision_payoff),
+             behavioral_belief_payoff = mean(behavioral_belief_payoff))
> behavioral_payoff

# A tibble: 4 × 3
  vis                behavioral_decision_payoff behavioral_belief_payoff
<chr>                <dbl>                <dbl>
1 HOPs                -6.81                -8.38
2 gradient            -6.59                -8.49
3 mean                -6.59                -8.65
4 mean+interval       -6.87                -8.56
```

Some results of the behavioral agents (-8.47 , -8.37 , -8.59 , and -8.33) are below the rational baseline (-7.96). Even though these results (in this example) come from some hypothetical behavioral agents who are simulated by quantal response equilibrium theory, this situation could also happen in a real experiment setting, where the behavioral agents might not follow a rational updating or decision rule or don't have a proper prior belief. Solely comparing these raw behavioral scores with the rational baseline and benchmark will entangle the loss caused by the information in visualization and the loss caused by the behavioral agents' improper prior belief or decision rule. Therefore, we propose the calibrated behavioral score, which can calibrate the behavioral agent by a Bayesian updating process and separate these two losses.

2.5.2 Calibrated behavioral score

We calibrate the behavioral agent by using a rational agent accessing their decisions as signals.

```
> # the joint distribution of actions and states
> joint_decision_distribution = data %>%
+   dplyr::group_by(mu, sigma, freezing_prob, behavioral_decision) %>%
+   summarise(count = n()) %>%
+   dplyr::group_by(mu, sigma, freezing_prob) %>%
+   mutate(count = count / sum(count))
> joint_belief_distribution = data %>%
+   dplyr::group_by(mu, sigma, freezing_prob, behavioral_belief) %>%
+   summarise(count = n()) %>%
+   dplyr::group_by(mu, sigma, freezing_prob) %>%
+   mutate(count = count / sum(count))
> sum_payoff = 0
> for (i in 1:nrow(data)) {
+   decision_signal = data[[i, "behavioral_decision"]]
+
+   joint_decision_distribution_filtered = joint_decision_distribution %>%
+     filter(behavioral_decision == decision_signal)
+
+   distributions = pnorm(0, mu, sigma_choices)
+
+   p_a = nrow(data %>% filter(behavioral_decision == decision_signal)) / nrow(data)
+
+   sigmas = joint_decision_distribution_filtered$sigma
+   p_a_d = as.list(rep(0, length(sigma_choices)))
+   names(p_a_d) = as.character(sigma_choices)
+   p_a_d[as.character(sigmas)] = joint_decision_distribution_filtered$count # p(a/d)
+   p_a_d = unname(unlist(p_a_d))
+   p_d_a = p_a_d * p_d / p_a # p(a/d) * p(d) / p(a)
+ }
```

```

+   calibrated_belief = sum(distributions * p_d_a) # E_p_d_a(d)
+   calibrated_reporting_decision = decision_space[[which.max(
+     lapply(decision_space, function(a) {
+       decision_scoring_rule(a, calibrated_belief)
+     })]]]
+   data[[i, "calibrated_decision_payoff"]] =
+     decision_scoring_rule(calibrated_reporting_decision, data[[i, "freezing_prob"]])
+
+   belief_signal = data[[i, "behavioral_belief"]]
+
+   joint_belief_distribution_filtered = joint_belief_distribution %>%
+     filter(behavioral_belief == belief_signal)
+
+   p_a = nrow(data %>% filter(behavioral_belief == belief_signal)) / nrow(data)
+
+   sigmas = joint_belief_distribution_filtered$sigma
+   p_a_d = as.list(rep(0, length(sigma_choices)))
+   names(p_a_d) = as.character(sigma_choices)
+   p_a_d[as.character(sigmas)] = joint_belief_distribution_filtered$count # p(a/d)
+   p_a_d = unname(unlist(p_a_d))
+   p_d_a = p_a_d * p_d / p_a # p(a/d) * p(d) / p(a)
+
+   calibrated_belief = sum(distributions * p_d_a) # E_p_d_a(d)
+   calibrated_reporting_belief = belief_space[[which.max(lapply(belief_space, function(a) {
+     belief_scoring_rule(a, calibrated_belief)
+   })]]]
+   data[[i, "calibrated_belief_payoff"]] =
+     belief_scoring_rule(calibrated_reporting_belief, data[[i, "freezing_prob"]])
+ }
> calibrated_payoff = data %>%
+   dplyr::group_by(vis) %>%
+   summarise(calibrated_decision_payoff = mean(calibrated_decision_payoff),
+             calibrated_belief_payoff = mean(calibrated_belief_payoff))
> calibrated_payoff

# A tibble: 4 × 3
  vis                calibrated_decision_payoff calibrated_belief_payoff
<chr>                <dbl>                <dbl>
1 HOPs                -6.81                -5.76
2 gradient            -6.59                -5.76
3 mean                -6.59                -5.77
4 mean+interval       -6.87                -5.77

```

Now we get the calibrated behavioral scores. The calibrated scores of all visualization types are between rational baseline and rational benchmark.